

## Howework 2: Classification of Cifar-10 images using Keras

Name: Akwasi Darkwah Akwaboah

Large Network (8 layers) 512 nodes each

```
#implementing in Keras
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras

#(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

#loading dataset
from keras.datasets import cifar10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

#1. Plotting the first 10 cifar10 images
# fig, axes = plt.subplots(1, 10, figsize=(10,5))
# for img, ax in zip(x_train[:10], axes):
#   ax.imshow(img)
# plt.show()

#2. Convert xtest and xtrain to float32 and normalize data to range 0.0 - 1.0
x_train_norm = np.float32(x_train/255.0)
x_test_norm = np.float32(x_test/255.0)
#print(x_train_norm)
#print(x_test_norm)

#3. function to convert to grayscale
def grayscale(data, dtype='float32'):
    # luma coding weighted average in video systems
    r, g, b = np.asarray(.3, dtype=dtype), np.asarray(.59, dtype=dtype), np.asarray(.11, dtype=dtype)
    rst = r * data[:, :, :, 0] + g * data[:, :, :, 1] + b * data[:, :, :, 2]
    # add channel dimension
    rst = np.expand_dims(rst, axis=3)
    return rst

#function Call
X_train_gray = grayscale(x_train_norm)
X_test_gray = grayscale(x_test_norm)

#plot a few grayscale vrs rgb
fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(x_test[:10], axes):
    ax.imshow(img)
plt.show()

fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(X_test_gray[:10], axes):
    ax.imshow(img[:, :, 0], cmap=plt.get_cmap('gray'), interpolation='none')
plt.show()

#create ANN model
#model constants

#define network
model = keras.Sequential([
    keras.layers.Flatten(input_shape = (32, 32, 1)),
    keras.layers.Dense(512, activation = 'relu'),
    keras.layers.Dense(512, activation = 'relu'),
    keras.layers.Dense(512, activation = 'relu'),
    keras.layers.Dense(512, activation = 'relu'),
```

```
keras.layers.Dense(512, activation = 'relu'),
keras.layers.Dense(512, activation = 'relu'),
keras.layers.Dense(512, activation = 'relu'),
keras.layers.Dense(512, activation = 'relu'),
keras.layers.Dense(10, activation = 'softmax')

])

#Compile Network
model.compile(optimizer='adam', loss ='sparse_categorical_crossentropy',
               metrics = ['accuracy'])

#train the network with 20% validation
history = model.fit(X_train_gray, y_train, batch_size = 128, epochs=20, validation_split=0.2) #spits

#test accuracy and loss
test_loss, test_acc = model.evaluate(X_test_gray, y_test, verbose=2)
print('\nTest accuracy:', test_acc)

predictions = model.predict(X_test_gray)

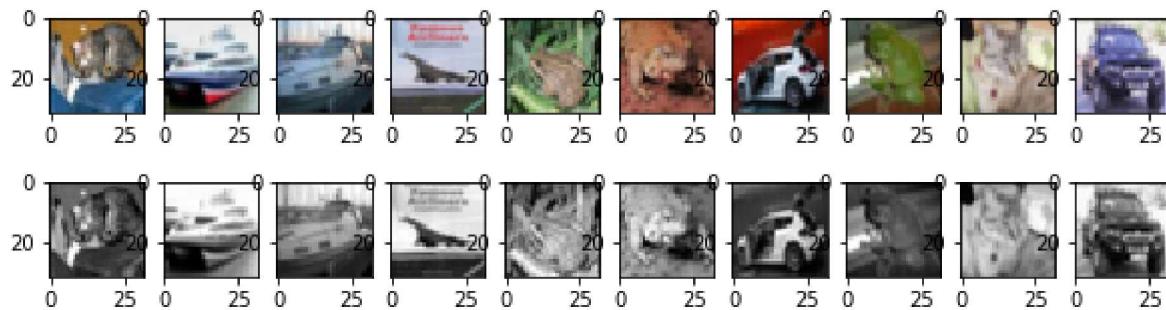
predictions[0]

np.argmax(predictions[0])

print(history.history.keys())
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.xlabel('Epochs')
plt.ylabel('Accuracy (fractional)')
plt.legend(['training accuracy', 'validation accuracy'], loc='best')

plt.figure()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.legend(['training loss', 'validation loss'], loc='best')
```





Train on 40000 samples, validate on 10000 samples

Epoch 1/20

40000/40000 [=====] - 19s 482us/sample - loss: 2.1207 - acc: 0.

Epoch 2/20

40000/40000 [=====] - 19s 472us/sample - loss: 1.9978 - acc: 0.

Epoch 3/20

40000/40000 [=====] - 19s 474us/sample - loss: 1.9430 - acc: 0.

Epoch 4/20

40000/40000 [=====] - 19s 469us/sample - loss: 1.8967 - acc: 0.

Epoch 5/20

40000/40000 [=====] - 19s 472us/sample - loss: 1.8584 - acc: 0.

Epoch 6/20

40000/40000 [=====] - 19s 473us/sample - loss: 1.8252 - acc: 0.

Epoch 7/20

40000/40000 [=====] - 19s 468us/sample - loss: 1.7967 - acc: 0.

Epoch 8/20

40000/40000 [=====] - 19s 470us/sample - loss: 1.7590 - acc: 0.

Epoch 9/20

40000/40000 [=====] - 19s 472us/sample - loss: 1.7303 - acc: 0.

Epoch 10/20

40000/40000 [=====] - 19s 468us/sample - loss: 1.7082 - acc: 0.

Epoch 11/20

40000/40000 [=====] - 19s 470us/sample - loss: 1.6792 - acc: 0.

Epoch 12/20

40000/40000 [=====] - 19s 480us/sample - loss: 1.6469 - acc: 0.

Epoch 13/20

40000/40000 [=====] - 19s 469us/sample - loss: 1.6201 - acc: 0.

Epoch 14/20

40000/40000 [=====] - 19s 468us/sample - loss: 1.5905 - acc: 0.

Epoch 15/20

40000/40000 [=====] - 19s 469us/sample - loss: 1.5645 - acc: 0.

Epoch 16/20

40000/40000 [=====] - 19s 469us/sample - loss: 1.5344 - acc: 0.

Epoch 17/20

40000/40000 [=====] - 19s 477us/sample - loss: 1.5048 - acc: 0.

Epoch 18/20

40000/40000 [=====] - 19s 470us/sample - loss: 1.4601 - acc: 0.

Epoch 19/20

40000/40000 [=====] - 19s 474us/sample - loss: 1.4345 - acc: 0.

Epoch 20/20

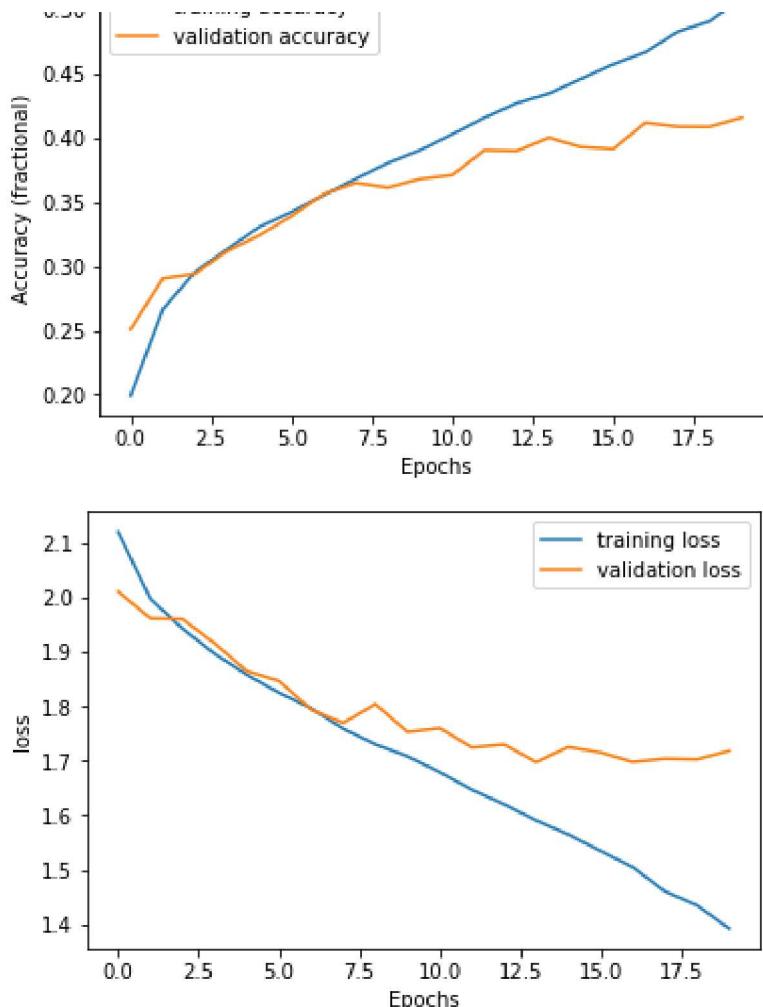
40000/40000 [=====] - 19s 472us/sample - loss: 1.3916 - acc: 0.

10000/10000 - 2s - loss: 1.7181 - acc: 0.4150

Test accuracy: 0.415

dict\_keys(['loss', 'acc', 'val\_loss', 'val\_acc'])

<matplotlib.legend at 0x7f781d9aab8>



Smaller network: 4-layers (tapering 512-256-64-32)

```
#implementing in Keras
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras

#(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

#loading dataset
from keras.datasets import cifar10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

#1. Plotting the first 10 cifar10 images
# fig, axes = plt.subplots(1, 10, figsize=(10,5))
# for img, ax in zip(x_train[:10], axes):
#     ax.imshow(img)
# plt.show()

#2. Convert xtest and xtrain to float32 and normalize data to range 0.0 - 1.0
x_train_norm = np.float32(x_train/255.0)
x_test_norm = np.float32(x_test/255.0)
#print(x_train_norm)
#print(x_test_norm)

#3. function to convert to grayscale
def grayscale(data, dtype='float32'):
    # luma coding weighted average in video systems
```

```

r, g, b = np.asarray(.3, dtype=dtype), np.asarray(.59, dtype=dtype), np.asarray(.11, dtype=dtype)
rst = r * data[:, :, 0] + g * data[:, :, 1] + b * data[:, :, 2]
# add channel dimension
rst = np.expand_dims(rst, axis=3)
return rst

#function Call
X_train_gray = grayscale(x_train_norm)
X_test_gray = grayscale(x_test_norm).

#plot a few grayscale vrs rgb
fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(x_test[:10], axes):
    ax.imshow(img)
plt.show()

fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(X_test_gray[:10], axes):
    ax.imshow(img[:, :, 0], cmap=plt.get_cmap('gray'), interpolation='none')
plt.show()

#create ANN model
#model constants

#define network
model = keras.Sequential([
    keras.layers.Flatten(input_shape = (32, 32, 1)),
    keras.layers.Dense(512, activation = 'relu'),
    keras.layers.Dense(256, activation = 'relu'),
    keras.layers.Dense(64, activation = 'relu'),
    keras.layers.Dense(32, activation = 'relu'),
    keras.layers.Dense(10, activation = 'softmax')
])
#Compile Network
model.compile(optimizer='adam', loss ='sparse_categorical_crossentropy',
               metrics = ['accuracy'])

#train the network with 20% validation
history = model.fit(X_train_gray, y_train, batch_size = 128, epochs=20, validation_split=0.2) #spits

#test accuracy and loss
test_loss, test_acc = model.evaluate(X_test_gray, y_test, verbose=2)
print('\nTest accuracy:', test_acc)

predictions = model.predict(X_test_gray)

predictions[0]

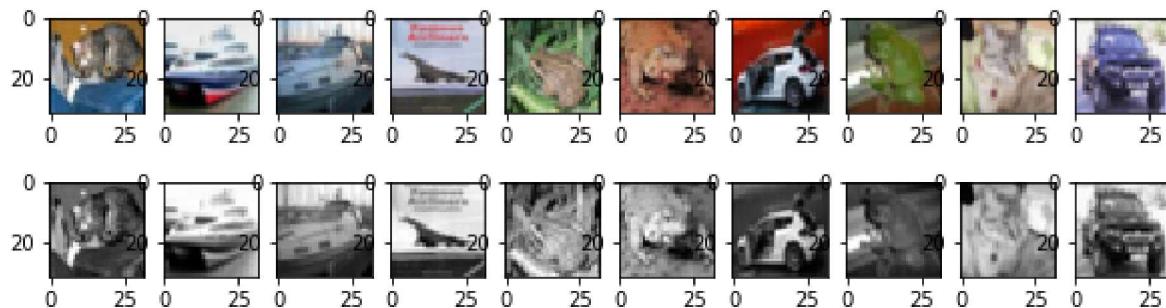
np.argmax(predictions[0])

print(history.history.keys())
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.xlabel('Epochs')
plt.ylabel('Accuracy (fractional)')
plt.legend(['training accuracy', 'validation accuracy'], loc='best')

plt.figure()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.legend(['training loss', 'validation loss'], loc='best')

```





Train on 40000 samples, validate on 10000 samples

Epoch 1/20

40000/40000 [=====] - 6s 138us/sample - loss: 2.0931 - acc: 0.2

Epoch 2/20

40000/40000 [=====] - 5s 127us/sample - loss: 1.9495 - acc: 0.2

Epoch 3/20

40000/40000 [=====] - 5s 126us/sample - loss: 1.8782 - acc: 0.3

Epoch 4/20

40000/40000 [=====] - 5s 127us/sample - loss: 1.8320 - acc: 0.3

Epoch 5/20

40000/40000 [=====] - 5s 124us/sample - loss: 1.7848 - acc: 0.3

Epoch 6/20

40000/40000 [=====] - 5s 123us/sample - loss: 1.7516 - acc: 0.3

Epoch 7/20

40000/40000 [=====] - 5s 122us/sample - loss: 1.7216 - acc: 0.3

Epoch 8/20

40000/40000 [=====] - 5s 123us/sample - loss: 1.6957 - acc: 0.3

Epoch 9/20

40000/40000 [=====] - 5s 122us/sample - loss: 1.6703 - acc: 0.4

Epoch 10/20

40000/40000 [=====] - 5s 123us/sample - loss: 1.6429 - acc: 0.4

Epoch 11/20

40000/40000 [=====] - 5s 122us/sample - loss: 1.6223 - acc: 0.4

Epoch 12/20

40000/40000 [=====] - 5s 122us/sample - loss: 1.5978 - acc: 0.4

Epoch 13/20

40000/40000 [=====] - 5s 122us/sample - loss: 1.5805 - acc: 0.4

Epoch 14/20

40000/40000 [=====] - 5s 123us/sample - loss: 1.5576 - acc: 0.4

Epoch 15/20

40000/40000 [=====] - 5s 121us/sample - loss: 1.5381 - acc: 0.4

Epoch 16/20

40000/40000 [=====] - 5s 123us/sample - loss: 1.5192 - acc: 0.4

Epoch 17/20

40000/40000 [=====] - 5s 126us/sample - loss: 1.5006 - acc: 0.4

Epoch 18/20

40000/40000 [=====] - 5s 126us/sample - loss: 1.4781 - acc: 0.4

Epoch 19/20

40000/40000 [=====] - 5s 124us/sample - loss: 1.4606 - acc: 0.4

Epoch 20/20

40000/40000 [=====] - 5s 123us/sample - loss: 1.4416 - acc: 0.4

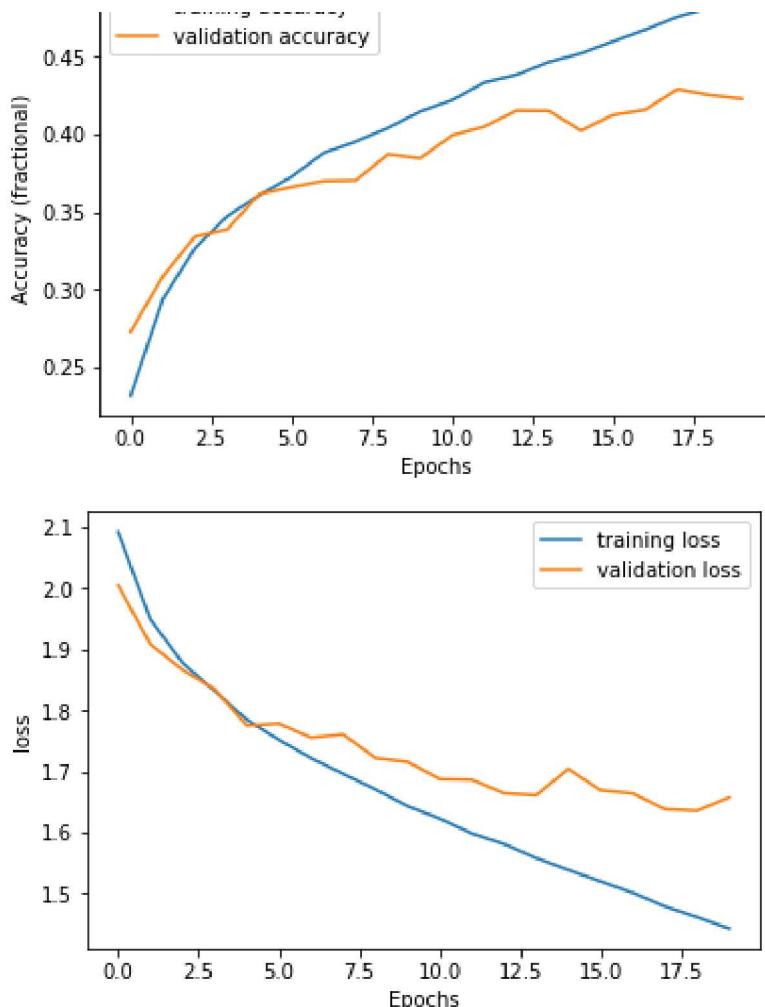
10000/10000 - 1s - loss: 1.6674 - acc: 0.4151

Test accuracy: 0.4151

dict\_keys(['loss', 'acc', 'val\_loss', 'val\_acc'])

<matplotlib.legend.Legend at 0x7f78170acef0>





## L2-norm regularization

```
#implementing in Keras
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras
from keras import regularizers

#(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

#loading dataset
from keras.datasets import cifar10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

#1. Plotting the first 10 cifar10 images
# fig, axes = plt.subplots(1, 10, figsize=(10,5))
# for img, ax in zip(x_train[:10], axes):
#   ax.imshow(img)
# plt.show()

#2. Convert xtest and xtrain to float32 and normalize data to range 0.0 - 1.0
x_train_norm = np.float32(x_train/255.0)
x_test_norm = np.float32(x_test/255.0)
#print(x_train_norm)
#print(x_test_norm)

#3. function to convert to grayscale
def grayscale(data, dtype='float32'):
```

```

# luma coding weighted average in video systems
r, g, b = np.asarray(.3, dtype=dtype), np.asarray(.59, dtype=dtype), np.asarray(.11, dtype=dtype)
rst = r * data[:, :, :, 0] + g * data[:, :, :, 1] + b * data[:, :, :, 2]
# add channel dimension
rst = np.expand_dims(rst, axis=3)
return rst

#function Call
X_train_gray = grayscale(x_train_norm)
X_test_gray = grayscale(x_test_norm)

#plot a few grayscale vrs rgb
fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(x_test[:10], axes):
    ax.imshow(img)
plt.show()

fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(X_test_gray[:10], axes):
    ax.imshow(img[:, :, 0], cmap=plt.get_cmap('gray'), interpolation='none')
plt.show()

#create ANN model
#model constants

#define network
l2_alpha = 0.0005

model = keras.Sequential([
    keras.layers.Flatten(input_shape = (32, 32, 1)),
    keras.layers.Dense(512, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(256, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(64, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(32, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(10, activation = 'softmax')
])
#Compile Network
model.compile(optimizer='adam', loss ='sparse_categorical_crossentropy',
              metrics = ['accuracy'])

#train the network with 20% validation
history = model.fit(X_train_gray, y_train, batch_size = 128, epochs=20, validation_split=0.2) #spits

#test accuracy and loss
test_loss, test_acc = model.evaluate(X_test_gray, y_test, verbose=2)
print('\nTest accuracy:', test_acc)

predictions = model.predict(X_test_gray)

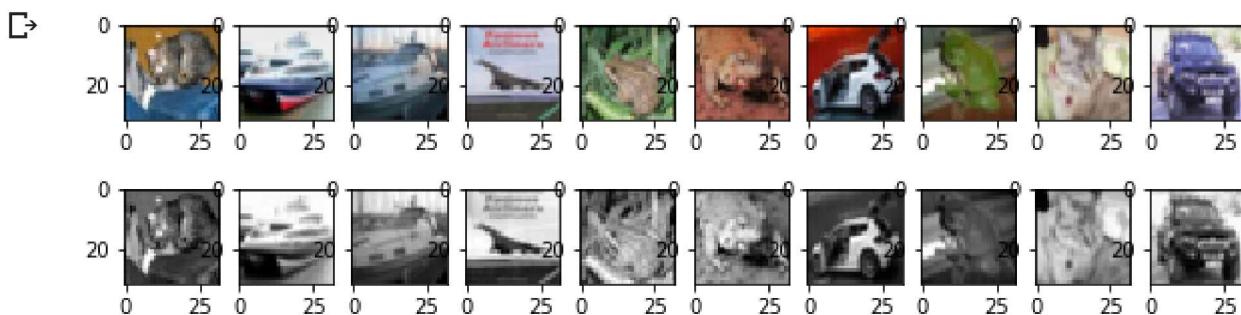
predictions[0]

np.argmax(predictions[0])

print(history.history.keys())
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.xlabel('Epochs')
plt.ylabel('Accuracy (fractional)')
plt.legend(['training accuracy', 'validation accuracy'], loc='best')

plt.figure()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.legend(['training loss', 'validation loss'], loc='best')

```



Train on 40000 samples, validate on 10000 samples

Epoch 1/20

40000/40000 [=====] - 6s 153us/sample - loss: 2.3419 - acc: 0.2

Epoch 2/20

40000/40000 [=====] - 5s 137us/sample - loss: 2.0735 - acc: 0.2

Epoch 3/20

40000/40000 [=====] - 5s 137us/sample - loss: 1.9953 - acc: 0.3

Epoch 4/20

40000/40000 [=====] - 6s 139us/sample - loss: 1.9417 - acc: 0.3

Epoch 5/20

40000/40000 [=====] - 6s 140us/sample - loss: 1.9130 - acc: 0.3

Epoch 6/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.8769 - acc: 0.3

Epoch 7/20

40000/40000 [=====] - 6s 141us/sample - loss: 1.8569 - acc: 0.3

Epoch 8/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.8384 - acc: 0.3

Epoch 9/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.8217 - acc: 0.3

Epoch 10/20

40000/40000 [=====] - 5s 137us/sample - loss: 1.8090 - acc: 0.3

Epoch 11/20

40000/40000 [=====] - 6s 141us/sample - loss: 1.7984 - acc: 0.3

Epoch 12/20

40000/40000 [=====] - 6s 140us/sample - loss: 1.7790 - acc: 0.3

Epoch 13/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.7676 - acc: 0.4

Epoch 14/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.7609 - acc: 0.4

Epoch 15/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.7450 - acc: 0.4

Epoch 16/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.7496 - acc: 0.4

Epoch 17/20

40000/40000 [=====] - 6s 141us/sample - loss: 1.7347 - acc: 0.4

Epoch 18/20

40000/40000 [=====] - 6s 139us/sample - loss: 1.7330 - acc: 0.4

Epoch 19/20

40000/40000 [=====] - 6s 138us/sample - loss: 1.7178 - acc: 0.4

Epoch 20/20

40000/40000 [=====] - 6s 140us/sample - loss: 1.7105 - acc: 0.4

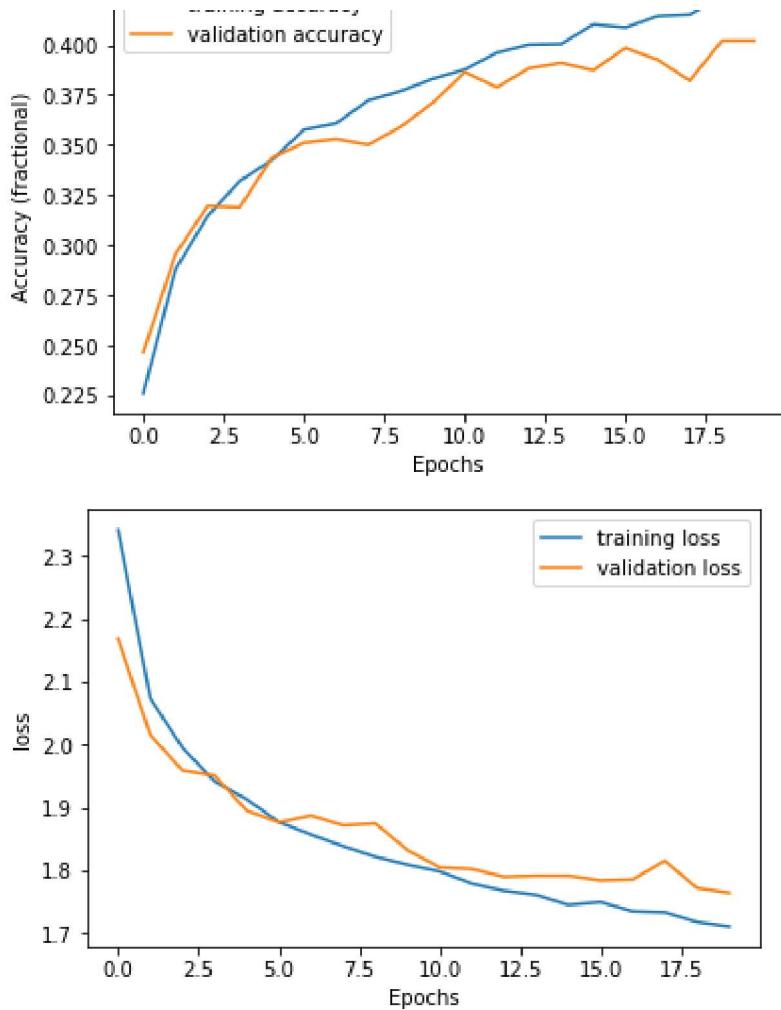
10000/10000 - 1s - loss: 1.7745 - acc: 0.3983

Test accuracy: 0.3983

```
dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
<matplotlib.legend.Legend at 0x7f7813816390>
```

0.425 | █ training accuracy

█



dropout

```
#implementing in Keras
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras
from keras import regularizers

#(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

#loading dataset
from keras.datasets import cifar10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

#1. Plotting the first 10 cifar10 images
# fig, axes = plt.subplots(1, 10, figsize=(10,5))
# for img, ax in zip(x_train[:10], axes):
#     ax.imshow(img)
# plt.show()

#2. Convert xtest and xtrain to float32 and normalize data to range 0.0 - 1.0
x_train_norm = np.float32(x_train/255.0)
x_test_norm = np.float32(x_test/255.0)
#print(x_train_norm)
#print(x_test_norm)

#3. function to convert to grayscale
def grayscale(data, dtype='float32'):
```

```

# luma coding weighted average in video systems
r, g, b = np.asarray(.3, dtype=dtype), np.asarray(.59, dtype=dtype), np.asarray(.11, dtype=dtype)
rst = r * data[:, :, 0] + g * data[:, :, 1] + b * data[:, :, 2]
# add channel dimension
rst = np.expand_dims(rst, axis=3)
return rst

#function Call
X_train_gray = grayscale(x_train_norm)
X_test_gray = grayscale(x_test_norm)

#plot a few grayscale vrs rgb
fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(x_test[:10], axes):
    ax.imshow(img)
plt.show()

fig, axes = plt.subplots(1, 10, figsize=(10,5))
for img, ax in zip(X_test_gray[:10], axes):
    ax.imshow(img[:, :, 0], cmap=plt.get_cmap('gray'), interpolation='none')
plt.show()

#create ANN model
#model constants

#define network
l2_alpha = 0.0005
drop_rate = 0.2

model = keras.Sequential([
    keras.layers.Flatten(input_shape = (32, 32, 1)),
    keras.layers.Dense(512, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(256, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(64, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dense(32, activation = 'relu', kernel_regularizer=regularizers.l2(l2_alpha)),
    keras.layers.Dropout(drop_rate, noise_shape=None, seed=None),
    keras.layers.Dense(10, activation = 'softmax'),
])
#Compile Network
model.compile(optimizer='adam', loss ='sparse_categorical_crossentropy',
              metrics = ['accuracy'])

#train the network with 20% validation
history = model.fit(X_train_gray, y_train, batch_size = 128, epochs=20, validation_split=0.2) #spits

#test accuracy and loss
test_loss, test_acc = model.evaluate(X_test_gray, y_test, verbose=2)
print('\nTest accuracy:', test_acc)

predictions = model.predict(X_test_gray)

predictions[0]

np.argmax(predictions[0])

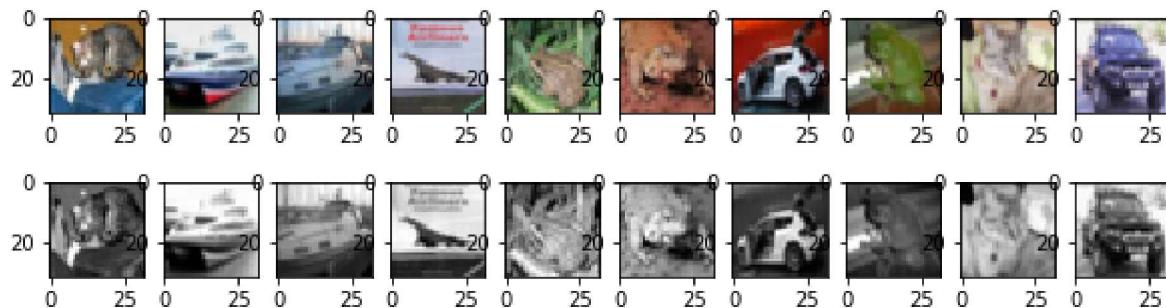
print(history.history.keys())
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.xlabel('Epochs')
plt.ylabel('Accuracy (fractional)')
plt.legend(['training accuracy', 'validation accuracy'], loc='best')

plt.figure()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.legend(['training loss', 'validation loss'], loc='best')

```

|





Train on 40000 samples, validate on 10000 samples

Epoch 1/20

40000/40000 [=====] - 5s 137us/sample - loss: 2.4127 - acc: 0.1

Epoch 2/20

40000/40000 [=====] - 5s 130us/sample - loss: 2.1266 - acc: 0.2

Epoch 3/20

40000/40000 [=====] - 5s 131us/sample - loss: 2.0335 - acc: 0.2

Epoch 4/20

40000/40000 [=====] - 5s 133us/sample - loss: 1.9858 - acc: 0.3

Epoch 5/20

40000/40000 [=====] - 5s 130us/sample - loss: 1.9552 - acc: 0.3

Epoch 6/20

40000/40000 [=====] - 5s 129us/sample - loss: 1.9191 - acc: 0.3

Epoch 7/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.8969 - acc: 0.3

Epoch 8/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.8890 - acc: 0.3

Epoch 9/20

40000/40000 [=====] - 5s 130us/sample - loss: 1.8668 - acc: 0.3

Epoch 10/20

40000/40000 [=====] - 5s 131us/sample - loss: 1.8497 - acc: 0.3

Epoch 11/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.8363 - acc: 0.3

Epoch 12/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.8269 - acc: 0.3

Epoch 13/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.8113 - acc: 0.3

Epoch 14/20

40000/40000 [=====] - 5s 134us/sample - loss: 1.8034 - acc: 0.3

Epoch 15/20

40000/40000 [=====] - 5s 133us/sample - loss: 1.7962 - acc: 0.3

Epoch 16/20

40000/40000 [=====] - 5s 135us/sample - loss: 1.7890 - acc: 0.3

Epoch 17/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.7757 - acc: 0.3

Epoch 18/20

40000/40000 [=====] - 5s 131us/sample - loss: 1.7683 - acc: 0.3

Epoch 19/20

40000/40000 [=====] - 5s 130us/sample - loss: 1.7577 - acc: 0.4

Epoch 20/20

40000/40000 [=====] - 5s 132us/sample - loss: 1.7544 - acc: 0.4

10000/10000 - 1s - loss: 1.8092 - acc: 0.3818

Test accuracy: 0.3818

dict\_keys(['loss', 'acc', 'val\_loss', 'val\_acc'])

<matplotlib.legend at 0x7f6f6c99ccf8>



